

4.1)

a) generalized\_policy\_iteration.m

```
% generalized_policy_iteration: Function solving the given MDP using the
%                               Generalized Policy Iteration algorithm
%
% Inputs:
%   world:           A structure defining the MDP to be solved
%   precision_pi:    Maximum value function change before
%                   terminating Policy Improvement step
%   max_ite_pi:       Maximum number of iterations for Policy
%                   Improvement loop
%   precision_pe:     Maximum value function change before
%                   terminating Policy Evaluation step
%   max_ite_pe:       Maximum number of iterations for Policy
%                   Evaluation loop
%
% Outputs:
%   V:               An array containing the value at each state
%   policy_index:    An array summarizing the index of the
%                   optimal action index at each state
%
% --
% Control for Robotics
% AER1517 Spring 2022
% Assignment 4
%
% --
% University of Toronto Institute for Aerospace Studies
% Dynamic Systems Lab
%
% Course Instructor:
% Angela Schoellig
% schoellig@utias.utoronto.ca
%
% Teaching Assistant:
% SiQi Zhou
% siqi.zhou@robotics.utias.utoronto.ca
% Lukas Brunke
% lukas.brunke@robotics.utias.utoronto.ca
% Adam Hall
% adam.hall@robotics.utias.utoronto.ca
%
% This script is adapted from the course on Optimal & Learning Control for
% Autonomous Robots at the Swiss Federal Institute of Technology in Zurich
% (ETH Zurich). Course Instructor: Jonas Buchli. Course Webpage:
% http://www.adrlab.org/doku.php/adrl:education:lecture:fs2015
%
% --
% Revision history
% [20.03.07, SZ]    first version

function [V, policy_index] = generalized_policy_iteration(world, precision_pi,
precision_pe, max_ite_pi, max_ite_pe)
    %% Initialization
    % MDP
    mdp = world.mdp;
    T = mdp.T;
    R = mdp.R;
    gamma = mdp.gamma;
    iteration_pi = 0;
    % Dimensions
    num_actions = length(T);
```

```

num_states = size(T{1}, 1);

% Initialize value function
V = zeros(num_states, 1);

% Initialize policy
% Note: Policy here encodes the action to be executed at state s. We
%       use deterministic policy here (e.g., [0,1,0,0] means take
%       action indexed 2)
random_act_index = randi(num_actions, [num_states, 1]);
policy = zeros(num_states, num_actions);
for s = 1:1:num_states
    selected_action = random_act_index(s);
    policy(s, selected_action) = 1;
end

while true
    %% [TODO] policy Evaluation (PE) (Section 2.6 of [1])
    iterations_pe = 0
    iteration_pi = iteration_pi + 1
    while iterations_pe <= max_ite_pe
        delta = 0
        iterations_pe = iterations_pe + 1
        for s = 1:1:num_states %loop for each state

            v = V(s,1); %initialize v value
            cur_state_index = s
            action_index = find(policy(s,:))
            noise_alpha = 0
            [next_state_index, next_state_noisy_index, reward] = ...
            one_step_gw_model(world, cur_state_index, action_index, noise_alpha)
%compute next state, reward of transition when applying a = pi(s)
            V(s,1) = reward + gamma*V(next_state_index,1) %compute value of
policy
            abs_diff = abs(v-V(s,1));
            delta = max(abs_diff,delta);
        end

        if delta < precision_pe

            break

        end
    end

    % V = ...;

    %% [TODO] Policy Improvement (PI) (Section 2.7 of [1])
    policyISstable = true;
    for s = 1:1:num_states

        b = policy(s,:);

        %compute argmax of cumulative reward function:
        cur_state_index = s;
        noise_alpha = 0;

        for a = 1:4
            [next_state_index, next_state_noisy_index, reward] =
one_step_gw_model(world, cur_state_index, a, noise_alpha)
            V_temp = reward + gamma*V(next_state_index,1);

            if a == 1
                temp = V_temp
            end
        end
    end
end

```

```

        end

        if V_temp >= temp
            temp = V_temp;
            argmax_a = a;    %computing which action maximizes cumulative
reward

        end

        end

        policy(s,:) = zeros(1,4);    %updating policy
        policy(s,argmax_a) = 1;

        if find(b) ~= find(policy(s,:))
            fprintf('policy not stable')
            policyISstable = false;
        end

    end

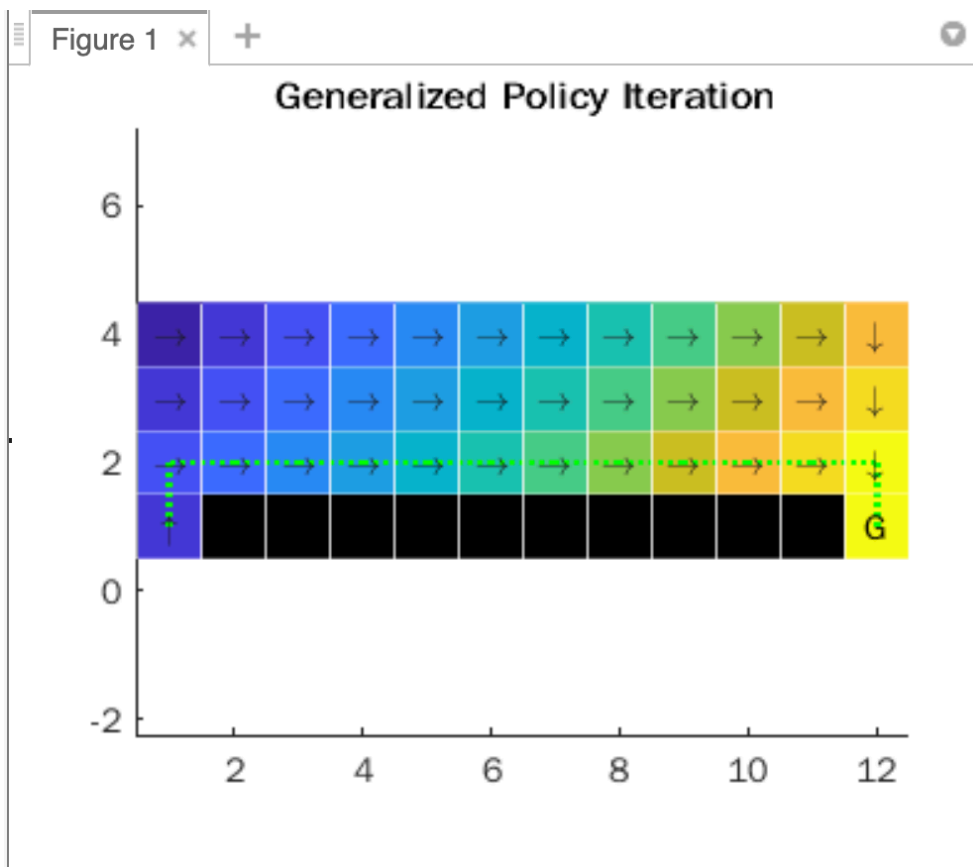
    if policyISstable == true

        break
    end
    % policy = ...;

    % Check algorithm convergence
    % if ...
    %     break
    % end
end
iteration_pi
% Return deterministic policy for plotting
[~, policy_index] = max(policy, [], 2);
end

```

## Heat Map and Results:



Difference between VI and PI:

In value iteration, the policy is evaluated only once and then improved.  
Value iteration converges faster than PI.

c) Monte Carlo

1) code: monte\_carlo.m:

```
% monte_carlo: Function solving the given MDP using the on-policy Monte
%               Carlo method
%
% Inputs:
%   world:       A structure defining the MDP to be solved
%   epsilon:     A parameter defining the 'softness' of the
%               epsilon-soft policy
%   k_epsilon:   The decay factor of epsilon per iteration
%   omega:       Learning rate for updating Q
%   training_itations: Maximum number of training episodes
%   episode_length: Maximum number of steps in each training
%               episodes
%
% Outputs:
```

```

%      Q:                An array containing the action value for
%                        each state-action pair
%      policy_index:      An array summarizing the index of the
%                        optimal action index at each state
%
%  --
%  Control for Robotics
%  AER1517 Spring 2022
%  Assignment 4
%
%  --
%  University of Toronto Institute for Aerospace Studies
%  Dynamic Systems Lab
%
%  Course Instructor:
%  Angela Schoellig
%  schoellig@utias.utoronto.ca
%
%  Teaching Assistant:
%  SiQi Zhou
%  siqi.zhou@robotics.utias.utoronto.ca
%  Lukas Brunke
%  lukas.brunke@robotics.utias.utoronto.ca
%  Adam Hall
%  adam.hall@robotics.utias.utoronto.ca
%
%  This script is adapted from the course on Optimal & Learning Control for
%  Autonomous Robots at the Swiss Federal Institute of Technology in Zurich
%  (ETH Zurich). Course Instructor: Jonas Buchli. Course Webpage:
%  http://www.adrlab.org/doku.php/adrl:education:lecture:fs2015
%
%  --
%  Revision history
%  [20.03.07, SZ]    first version

function [Q, policy_index] = ...
    monte_carlo(world, epsilon, k_epsilon, omega, training_iterations,
episode_length)
    %% Initialization
    % MDP
    mdp = world.mdp;
    gamma = mdp.gamma;

    % States
    STATES = mdp.STATES;
    ACTIONS = mdp.ACTIONS;

    % Dimensions
    num_states = size(STATES, 2);
    num_actions = size(ACTIONS, 2);

    % Create object for incremental plotting of reward after each episode
    windowSize = 10; %Sets the width of the sliding window filter used in
plotting
    plotter = RewardPlotter(windowSize);

    % Initialize Q
    Q = zeros(num_states, num_actions);
    terminal_state = 12
    % [TODO] Initialize epsilon-soft policy
    % policy = ...; % size: num_states x num_actions
    policy = zeros(num_states, num_actions);
    policy = initialize_random_policy(epsilon, num_states, num_actions);
    %% On-policy Monte Carlo Algorithm (Section 2.9.3 of [1])
    initial_state = randi([1, num_states]);
    curr_state = initial_state;

```

```

cur_state_index = curr_state;
state_sequence = [curr_state];
reward_sequence = [];
action_sequence = [];

for train_loop = 1:1:training_iterations
    %% [TODO] Generate a training episode
    initial_state = randi([1, num_states]); %randomly initialize states
    cur_state_index = initial_state;
    R = 0; %Initialize Return
    state_sequence = [curr_state];
    reward_sequence = []; %initialize reward sequence
    action_sequence = []; %initialize state sequence
    episode_index = 0
    while cur_state_index ~= terminal_state & episode_index <
episode_length % episode termination criteria
        episode_index = episode_index + 1;
        policy_prob = policy(cur_state_index,:);
        % Sample current epsilon-soft policy
        action = sample_from_epsilon_policy(epsilon,policy_prob);

        % Interaction with environment
        [next_state_index, ~, reward] = one_step_gw_model(world,
cur_state_index, action, 1);
        state_sequence = [state_sequence,next_state_index];
        reward_sequence = [reward_sequence, reward];
        action_sequence = [action_sequence, action];

        cur_state_index = next_state_index;
        % Log data for the episode
        % ...
    end
    N = length(state_sequence);
    i = 0;
    reward_sequence
    action_sequence
    state_sequence
    for i = 1:N-1

        s = state_sequence(N-i);
        a = action_sequence(N-i);
        r = reward_sequence(N-i)
        R = gamma*R + r ; % cumulative return
        Q(s,a) = Q(s,a) + omega*(R - Q(s,a));

    end
    R
    % Update Q(s,a)
    % Q = ...;

    %% [TODO] Update policy(s,a)
    for i = 1:N-1

        x = state_sequence(N-i);
        u_optim = arg_max_Q(Q,x);

        for a = 1:4

            if a == u_optim
                policy(x,u_optim) = 1 - 0.75*epsilon;
            else
                policy(x,a) = 0.25*epsilon;
            end
        end
    end
end

```

```

    end

    end

    % policy = ...;


    %% [TODO] Update the reward plot
    EpisodeTotalReturn = R; % Sum of the reward obtained during the episode
    plotter = UpdatePlot(plotter, EpisodeTotalReturn);
    drawnow;
    pause(0.1);


    %% Decrease the exploration
    % Set k_epsilon = 1 to maintain constant exploration
    epsilon = epsilon * k_epsilon;
end

% Return deterministic policy for plotting
[~, policy_index] = max(policy, [], 2);
end

```

### Results (Monte Carlo)





2. Impact of varying policy parameter epsilon:



d)

Code: q\_learning.m:

```
% q_learning: Function solving the given MDP using the off-policy
%               Q-Learning method
%
% Inputs:
%   world:      A structure defining the MDP to be solved
%   epsilon:    A parameter defining the 'softness' of the
%               epsilon-soft policy
%   k_epsilon:  The decay factor of epsilon per iteration
%   omega:      Learning rate for updating Q
%   training_iterations: Maximum number of training episodes
%   episode_length: Maximum number of steps in each training
%                   episodes
%   noise_alpha: A parameter that controls the noisiness of
%               observation (observation is noise-free when
%               noise_alpha is set to 1 and is more
%               corrupted when it is set to values closer
%               to 0)
%
% Outputs:
%   Q:          An array containing the action value for
%               each state-action pair
%   policy_index: An array summarizing the index of the
%               optimal action index at each state
%
% --
% Control for Robotics
% AER1517 Spring 2022
% Assignment 4
%
% --
% University of Toronto Institute for Aerospace Studies
% Dynamic Systems Lab
%
% Course Instructor:
% Angela Schoellig
% schoellig@utias.utoronto.ca
%
% Teaching Assistant:
% SiQi Zhou
% siqi.zhou@robotics.utias.utoronto.ca
% Lukas Brunke
% lukas.brunke@robotics.utias.utoronto.ca
% Adam Hall
% adam.hall@robotics.utias.utoronto.ca
%
% This script is adapted from the course on Optimal & Learning Control for
% Autonomous Robots at the Swiss Federal Institute of Technology in Zurich
% (ETH Zurich). Course Instructor: Jonas Buchli. Course Webpage:
% http://www.adrlab.org/doku.php/adrl:education:lecture:fs2015
%
% --
% Revision history
% [20.03.07, SZ]    first version

function [Q, policy_index] = q_learning(world, epsilon, k_epsilon, omega,
training_iterations, episode_length, noise_alpha)
    %% Initialization
    % MDP
    mdp = world.mdp;
    gamma = mdp.gamma;
    terminal_state = 12
```

```

% States
STATES = mdp.STATES;
ACTIONS = mdp.ACTIONS;

% Dimensionts
num_states = size(STATES, 2);
num_actions = size(ACTIONS, 2);

% Create object for incremental plotting of reward after each episode
windowSize = 10; %Sets the width of the sliding window fitler used in
plotting
plotter = RewardPlotter(windowSize);

% Initialize Q
Q = zeros(num_states, num_actions);

% [TODO] Initialize epsilon-soft policy
% policy = ...; % size: num_states x num_actions
policy = initialize_random_policy(epsilon,num_states,num_actions);
%% Q-Learning Algorithim (Section 2.9 of [1])
for train_loop = 1:1:training_iterations
    %% [TODO] Generate a training episode
    initial_state = randi([1, num_states]); %randomly initialize states
%     initial_state = mdp.s_start_index;
    cur_state_index = initial_state;
    reward_sequence = [];
    episode_index = 0;
    R = 0;
    while cur_state_index ~= terminal_state & episode_index <=
episode_length
        % Sample current epsilon-soft policy
        policy_prob = policy(cur_state_index,:);
        % Sample current epsilon-soft policy
        action =sample_from_epsilon_policy(epsilon,policy_prob);
        episode_index = episode_index + 1;
        % Interaction with environment
        % Note: 'next_state_noisy_index' below simulates state
        %       observarions corrupted with noise. Use this for
        %       Q-learning correspondingly for the last part of
        %       Problem 2.2 (d)
        [next_state_index, next_state_noisy_index, reward] = ...
            one_step_gw_model(world, cur_state_index, action, noise_alpha);
        %Q learning update rule::
        argmax_u_prime = arg_max_Q(Q, next_state_index);
        Q(cur_state_index,action) = Q(cur_state_index,action) +
omega*(reward + gamma*Q(next_state_index,argmax_u_prime) -
Q(cur_state_index,action));
        reward_sequence = [reward_sequence, reward];

        % Log data for the episode
        % ...
        x = cur_state_index;
        u_optim = arg_max_Q(Q,x);

        for a = 1:4

            if a == u_optim
                policy(x,u_optim) = 1 - 0.75*epsilon;
            else
                policy(x,a) = 0.25*epsilon;
            end

            % Update Q(s,a)
            % Q = ...;
        end
    end
end

```

```

%% [TODO] Update policy(s,a)
% policy = ...;
N = length(reward_sequence);
for i = 1:N-1

    r = reward_sequence(N-i+1);
    R = gamma*R + r ; % cumulative return

end

cur_state_index = next_state_index;
end

%% [TODO] Update the reward plot
EpisodeTotalReturn = R %Sum of the reward obtained during the episode
plotter = UpdatePlot(plotter, EpisodeTotalReturn);
drawnow;
pause(0.1);


%% Decrease the exploration
k_epsilon = 1 %to maintain constant exploration
epsilon = epsilon * k_epsilon;


end

% Return deterministic policy for plotting
[~, policy_index] = max(policy, [], 2);
end

```

Simulation Results:





Impact of changing decay parameter  $k\_epsilon$

Difference between monte carlo and q learning algorithm:

Q learning with noise:



4.2)

4.2.a)

1.1 Code - build\_stochastic\_mdp\_nn.m

```
% build_stochastic_mdp_nn: Function implementing the Nearest Neighbour
%                               approach for creating a stochastic MDP
%
% Inputs:
%   world:      A structure containing basic parameters for
%               the mountain car problem
%   T:          Transition model with elements initialized
%               to zero
%   R:          Expected reward model with elements
%               initialized to zero
%   num_samples: Number of samples to use for creating the
%                stochastic model
%
% Outputs:
%   T:          Transition model with elements  $T\{a\}(s,s')$ 
%               being the probability of transition to
%               state  $s'$  from state  $s$  taking action  $a$ 
%   R:          Expected reward model with elements
%                $R\{a\}(s,s')$  being the expected reward on
```

```

%                                     transition from s to s' under action a
%
% --
% Control for Robotics
% AER1517 Spring 2022
% Assignment 4
%
% --
% University of Toronto Institute for Aerospace Studies
% Dynamic Systems Lab
%
% Course Instructor:
% Angela Schoellig
% schoellig@utias.utoronto.ca
%
% Teaching Assistant:
% SiQi Zhou
% siqi.zhou@robotics.utias.utoronto.ca
% Lukas Brunke
% lukas.brunke@robotics.utias.utoronto.ca
% Adam Hall
% adam.hall@robotics.utias.utoronto.ca
%
% --
% Revision history
% [20.03.07, SZ]    first version

function [T, R] = build_stochastic_mdp_nn(world, T, R, num_samples)
    % Extract states and actions
    STATES = world.mdp.STATES;
    ACTIONS = world.mdp.ACTIONS;

    % Dimensions
    num_states = size(STATES, 2);
    num_actions = size(ACTIONS, 2);

    % Loop through all possible states
    for state_index = 1:1:num_states
        cur_state = STATES(:, state_index);
        fprintf('building model... state %d\n', state_index);

        % Apply each possible action
        for action_index = 1:1:num_actions
            action = ACTIONS(:, action_index);
            %
            p_k = cur_state(1);
            %
            v_k = cur_state(2);
            %
            v_k_next = v_k + 0.001*action_index

            % [TODO] Build a stochastic MDP based on Nearest Neighbour
            % Note: The function 'nearest_state_index_lookup' can be used
            % to find the nearest node to a continuous state
            for samples = 1:1:num_samples
                [next_state, reward, is_goal_state] = one_step_mc_model(world,
cur_state, action)
                next_state(1) = next_state(1) + normrnd(0,0.001);
                next_state(2) = next_state(2) + normrnd(0,0.005);
                next_state_nearest = nearest_state_index_lookup(STATES,
next_state);
                T{action_index}(state_index,next_state_nearest) =
T{action_index}(state_index,next_state_nearest) + 1/num_samples;
                % Update transition and reward models
                % T{action_index}(state_index, next_state_index) = ...;
                R{action_index}(state_index, next_state_nearest) = reward;
            end
        end
    end
end

```



```
        end
    end
end
```

1.2 What is the stochastic element in the modelling process and what is its significance ?

1.3 What modelling parameters would have the most impact on the quality of the solution ?

## 2.1) main\_p2\_mc\_rl.m

```
% main_p2_mc_rl: Main script for Problem 4.2 mountain car (RL approach)
%
% --
% Control for Robotics
% AER1517 Spring 2022
% Assignment 4
%
% --
% University of Toronto Institute for Aerospace Studies
% Dynamic Systems Lab
%
% Course Instructor:
% Angela Schoellig
% schoellig@utias.utoronto.ca
%
% Teaching Assistant:
% SiQi Zhou
% siqi.zhou@robotics.utias.utoronto.ca
% Lukas Brunke
% lukas.brunke@robotics.utias.utoronto.ca
% Adam Hall
% adam.hall@robotics.utias.utoronto.ca
%
% --
% Revision history
% [20.03.07, SZ]    first version

clear all;
close all;
clc;

%% General
% Add path
addpath(genpath(pwd));

% Result and plot directory
save_dir = './results/';
% mkdir(save_dir);

%% Problem 4.2 (a)–(b) Create stochastic MDPs for the mountain car problem
% [TODO] Load mountain car model
% change model name correspondingly:
%     (a) 'mountain_car_nn' for the nearest neighbour method
%     (b) 'mountain_car_li' for the linear interpolation approach
load('mountain_car_model/mountain_car_nn');

%% Generalized policy iteration
% Algorithm parameters
precision_pi = 0.1;
precision_pe = 0.01;
max_ite_pi = 100;
```

```

max_ite_pe = 10;

% Solve MDP
[v_gpi, policy_gpi] = generalized_policy_iteration_mc(world, precision_pi, ...
    precision_pe, max_ite_pi, max_ite_pe);

% Visualization
plot_value = true;
plot_flowfield = true;
plot_visualize = true;
plot_title = 'Generalized Policy Iteration';
hdl_gpi = visualize_mc_solution(world, v_gpi, policy_gpi, plot_value, ...
    plot_flowfield, plot_visualize, plot_title, save_dir);

% Save results
save(strcat(save_dir, 'gpi_results.mat'), 'v_gpi', 'policy_gpi');

```

### generalized\_policy\_iteration\_mc.m

```

% generalized_policy_iteration: Function solving the given MDP using the
%                               Generalized Policy Iteration algorithm
%
% Inputs:
%   world:           A structure defining the MDP to be solved
%   precision_pi:    Maximum value function change before
%                   terminating Policy Improvement step
%   max_ite_pi:      Maximum number of iterations for Policy
%                   Improvement loop
%   precision_pe:    Maximum value function change before
%                   terminating Policy Evaluation step
%   max_ite_pe:      Maximum number of iterations for Policy
%                   Evaluation loop
%
% Outputs:
%   V:               An array containing the value at each state
%   policy_index:    An array summarizing the index of the
%                   optimal action index at each state
%
% --
% Control for Robotics
% AER1517 Spring 2022
% Assignment 4
%
% --
% University of Toronto Institute for Aerospace Studies
% Dynamic Systems Lab
%
% Course Instructor:
% Angela Schoellig
% schoellig@utias.utoronto.ca
%
% Teaching Assistant:
% SiQi Zhou
% siqi.zhou@robotics.utias.utoronto.ca
% Lukas Brunke
% lukas.brunke@robotics.utias.utoronto.ca
% Adam Hall
% adam.hall@robotics.utias.utoronto.ca
%
% This script is adapted from the course on Optimal & Learning Control for
% Autonomous Robots at the Swiss Federal Institute of Technology in Zurich
% (ETH Zurich). Course Instructor: Jonas Buchli. Course Webpage:
% http://www.adrlab.org/doku.php/adrl:education:lecture:fs2015
%

```

```

% --
% Revision history
% [20.03.07, SZ]    first version

function [V, policy_index] = generalized_policy_iteration_mc(world,
precision_pi, precision_pe, max_ite_pi, max_ite_pe)
    %% Initialization
    % MDP
    mdp = world.mdp;
    T = mdp.T;
    R = mdp.R;
    gamma = mdp.gamma;

    % Discrete states
    POS = world.mdp.POS;
    VEL = world.mdp.VEL;

    % Dimensions
    num_actions = length(T);
    num_states = size(T{1}, 1);

    % Intialize value function
    V = zeros(num_states, 1);

    % Initialize policy
    % Note: Policy here encodes the action to be executed at state s. We
    %       use deterministic policy here (e.g., [0,1,0,0,0] means take
    %       action indexed 2)
    random_act_index = randi(num_actions, [num_states, 1]);
    policy = zeros(num_states, num_actions);
    for s = 1:1:num_states
        selected_action = random_act_index(s);
        policy(s, selected_action) = 1;
    end
    iterations_pi = 0;
    while iterations_pi <= max_ite_pi
        iterations_pe = 0;
        iterations_pi = iterations_pi + 1;
        %% [TODO] policy Evaluation (PE) (Section 2.6 of [1])
        while iterations_pe <= max_ite_pe
            delta = 0;
            iterations_pe = iterations_pe + 1;
            for s = 1:1:num_states
                v = V(s,1);

                %%%%Computation of V%%%%%%%%%
                temp_v = 0; %temporary variable for value function computation

                for a = 1:num_actions
                    temp_R = 0; %temporary variable for expected return
                    computation
                    for s_prime = 1:num_states
                        temp_R = temp_R + T{a}(s,s_prime)*(R{a}(s,s_prime) +
gamma*V(s_prime,1));
                    end
                    temp_v = temp_v + policy(s,a)*temp_R;
                end
                %finished computation of value function
                V(s,1) = temp_v;
                abs_diff = abs(v-V(s,1));
                delta = max(abs_diff,delta);
                if delta< precision_pe
                    break
                end
            end
        end
    end
end

```

```

end
% V = ...;

%% [TODO] Policy Improvement (PI) (Section 2.7 of [1])

policy_is_stable = true;

for s = 1:num_states
    b = policy(s,:);
    %compute argmax of value function
    for a = 1:num_actions
        temp_R = 0; %temporary variable for expected return computation
        for s_prime = 1:num_states
            temp_R = temp_R + T{a}(s,s_prime)*( R{a}(s,s_prime) +
gamma*V(s_prime,1));
        end
        if a==1
            temp = temp_R;
            arg_max = a;
        end

        if temp_R >= temp
            temp = temp_R;
            arg_max = a;
        end
    end

    policy(s,:) = zeros(1,num_actions); %updating policy
    policy(s,arg_max) = 1;

    if find(b) ~= find(policy(s,:))
        fprintf('policy not stable')
        policyISstable = false;
    end
end

if policyISstable == true
    break
end
% policy = ...;

% Check algorithm convergence
% if ...
%     break
% end
end

% Return deterministic policy for plotting
[~, policy_index] = max(policy, [], 2);
end



```


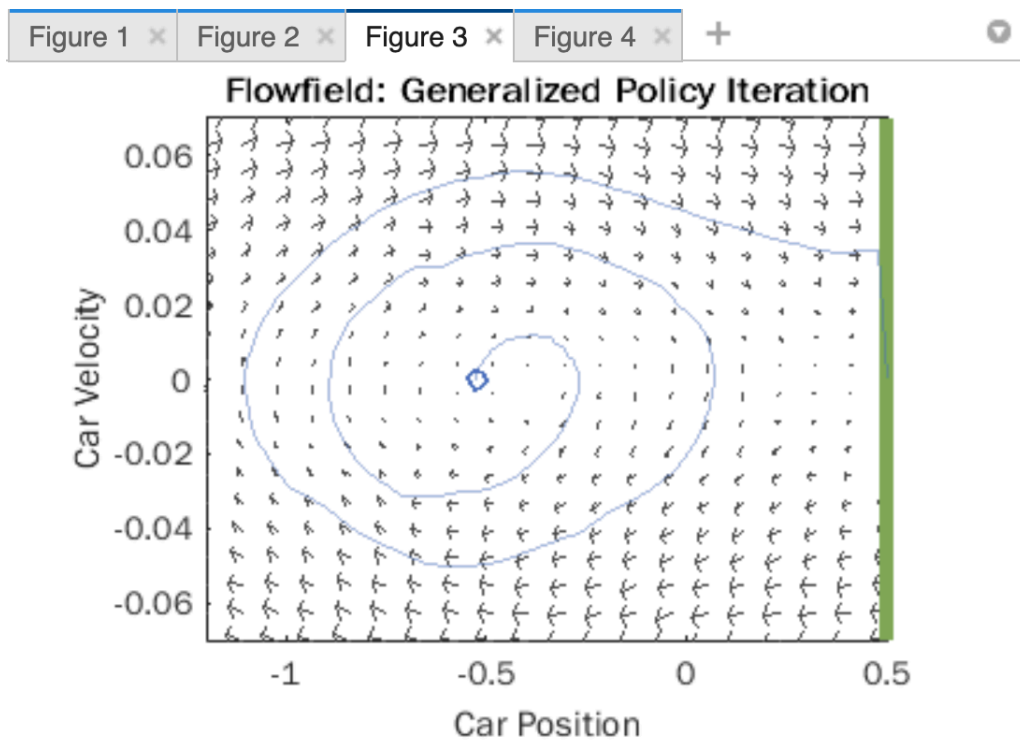
Q: Was the learning algorithm able to find this solution? If not, why do you think that is the case?

Yes.



## Converged Heat Map & Results(Nearest Neighbour Case)





```

% build_stochastic_mdp_li: Function implementing the Linear Interpolation
%                           approach for creating a stochastic MDP
%
% Inputs:
%   world:                  A structure containing basic parameters for
%                           the mountain car problem
%   T:                      Transition model with elements initialized
%                           to zero
%   R:                      Expected reward model with elements
%                           initialized to zero
%
% Outputs:
%   T:                      Transition model with elements  $T\{a\}(s,s')$ 
%                           being the probability of transition to
%                           state  $s'$  from state  $s$  taking action  $a$ 
%   R:                      Expected reward model with elements
%                            $R\{a\}(s,s')$  being the expected reward on
%                           transition from  $s$  to  $s'$  under action  $a$ 
%
% --
% Control for Robotics
% AER1517 Spring 2022
% Assignment 4
%
% --
% University of Toronto Institute for Aerospace Studies
% Dynamic Systems Lab
%
% Course Instructor:
% Angela Schoellig
% schoellig@utias.utoronto.ca
%
% Teaching Assistant:
% SiQi Zhou
% siqi.zhou@robotics.utias.utoronto.ca
% Lukas Brunke
% lukas.brunke@robotics.utias.utoronto.ca
% Adam Hall
% adam.hall@robotics.utias.utoronto.ca
%
% --
% Revision history
% [20.03.07, SZ]    first version

function [T, R] = build_stochastic_mdp_li(world, T, R)
    % Extract states and actions
    STATES = world.mdp.STATES;
    ACTIONS = world.mdp.ACTIONS;

    % Number of discrete states and actions
    num_states = size(STATES, 2);
    num_actions = size(ACTIONS, 2);

    % State space dimension
    dim_state = size(STATES, 1);

    % Unique values
    for i = 1:1:dim_state
        unique_states{i} = unique(STATES(i,:));
    end

    % Loop through all possible states
    for state_index = 1:1:num_states
        cur_state = STATES(:, state_index);
        fprintf('building model... state %d\n', state_index);
    end

```



```

% Apply each possible action
for action_index = 1:1:num_actions
    action = ACTIONS(:, action_index);

    % Propagate forward
    [next_state, reward, ~] = world.one_step_model(world, ...
        cur_state, action);

    % Find four vertices enclosing next state index
    for i = 1:1:dim_state
        % find closest discretized values along state dimension i
        node_index_temp = knnsearch(unique_states{i}', next_state(i),
'K', 2);
        node_value_temp = unique_states{i}(node_index_temp);

        % for each state dimension i, store the min-max bounds
        box_min = min(node_value_temp);
        box_max = max(node_value_temp);
        node_value(i,1:2) = [box_min, box_max];

        % normalize next state values
        next_state_normalized(i,1) = ...
            (next_state(i,1) - box_min) / (box_max - box_min);
    end

    % node values (for two-dim state space)
    node(1:2,1) = [node_value(1,1); node_value(2,1)]; % lower-left
    node(1:2,2) = [node_value(1,2); node_value(2,1)]; % lower-right
    node(1:2,3) = [node_value(1,2); node_value(2,2)]; % upper-right
    node(1:2,4) = [node_value(1,1); node_value(2,2)]; % upper-left

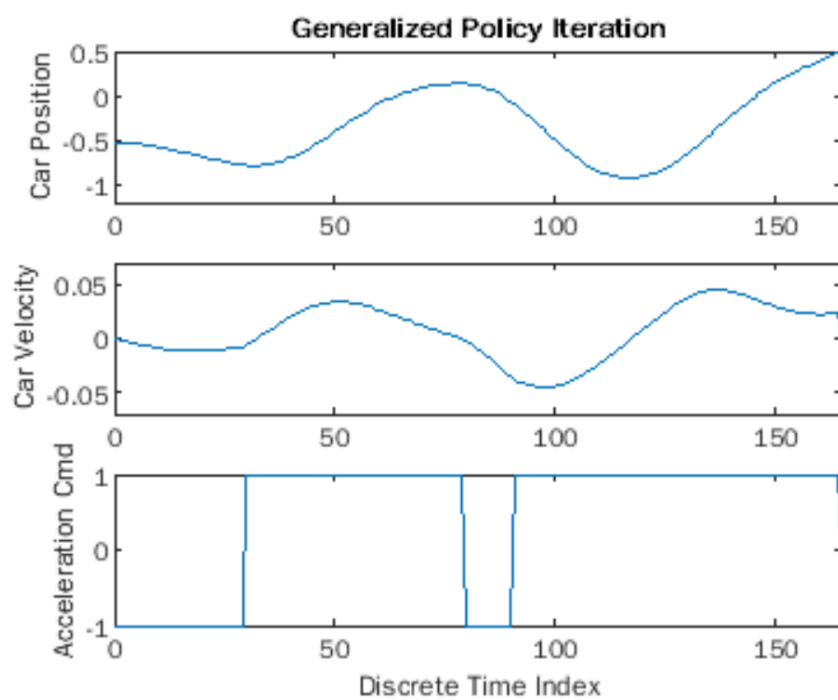
    % [TODO] Assign probability to adjacent nodes (bilinear)
    x = next_state_normalized(1);
    y = next_state_normalized(2);
    prob(1) = (1-x)*(1-y); % min min
    prob(2) = x*(1-y); % max min
    prob(3) = x*y; % max max
    prob(4) = (1-x)*(y); % min max

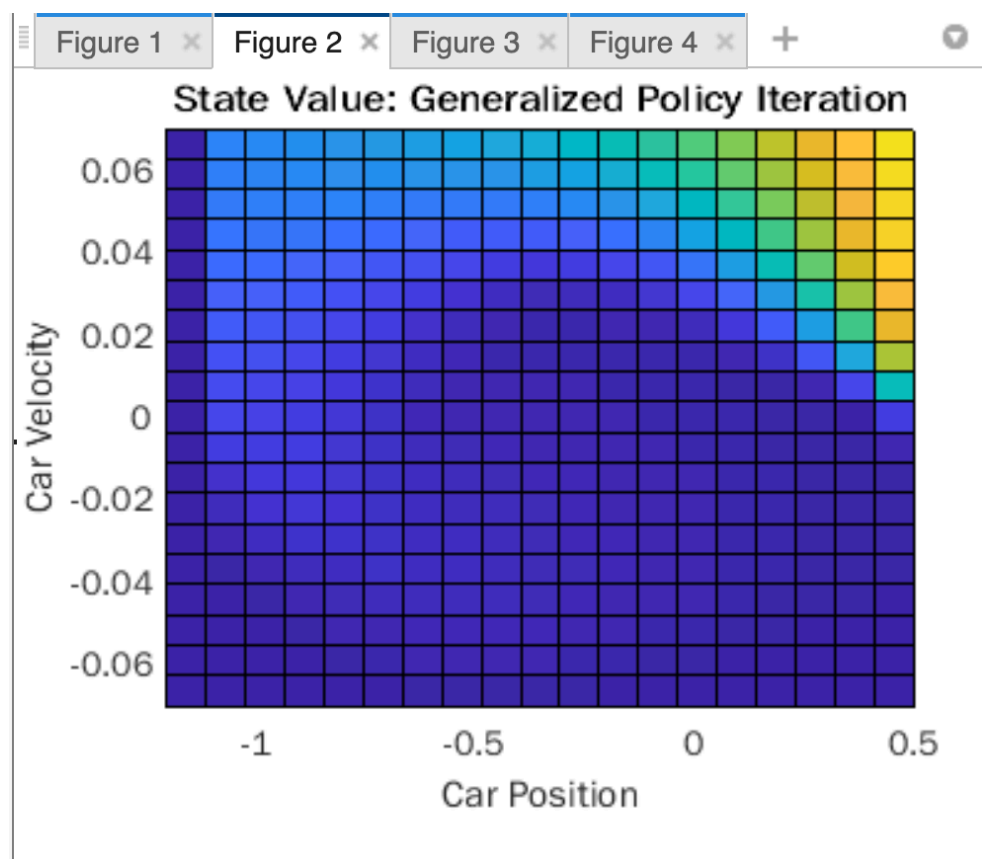
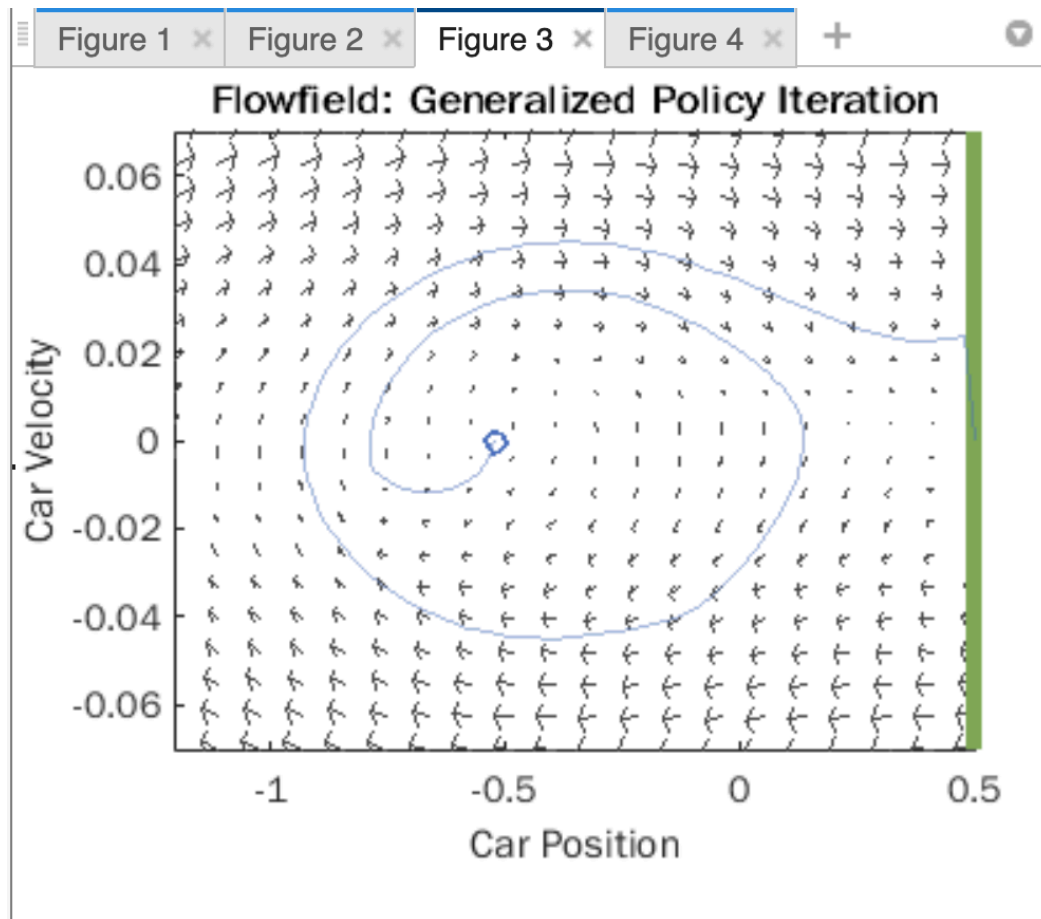
    % Update probability and reward for each node
    for i = 1:1:4
        node_index = nearest_state_index_lookup(STATES, node(:,i));


        % Update transition and reward models
        T{action_index}(state_index, node_index) = prob(i);
        R{action_index}(state_index, node_index) = reward;
    end
end
end
end
end

```

Figure 1 × Figure 2 × Figure 3 × Figure 4 × +







Effectiveness of LI and NN:

